
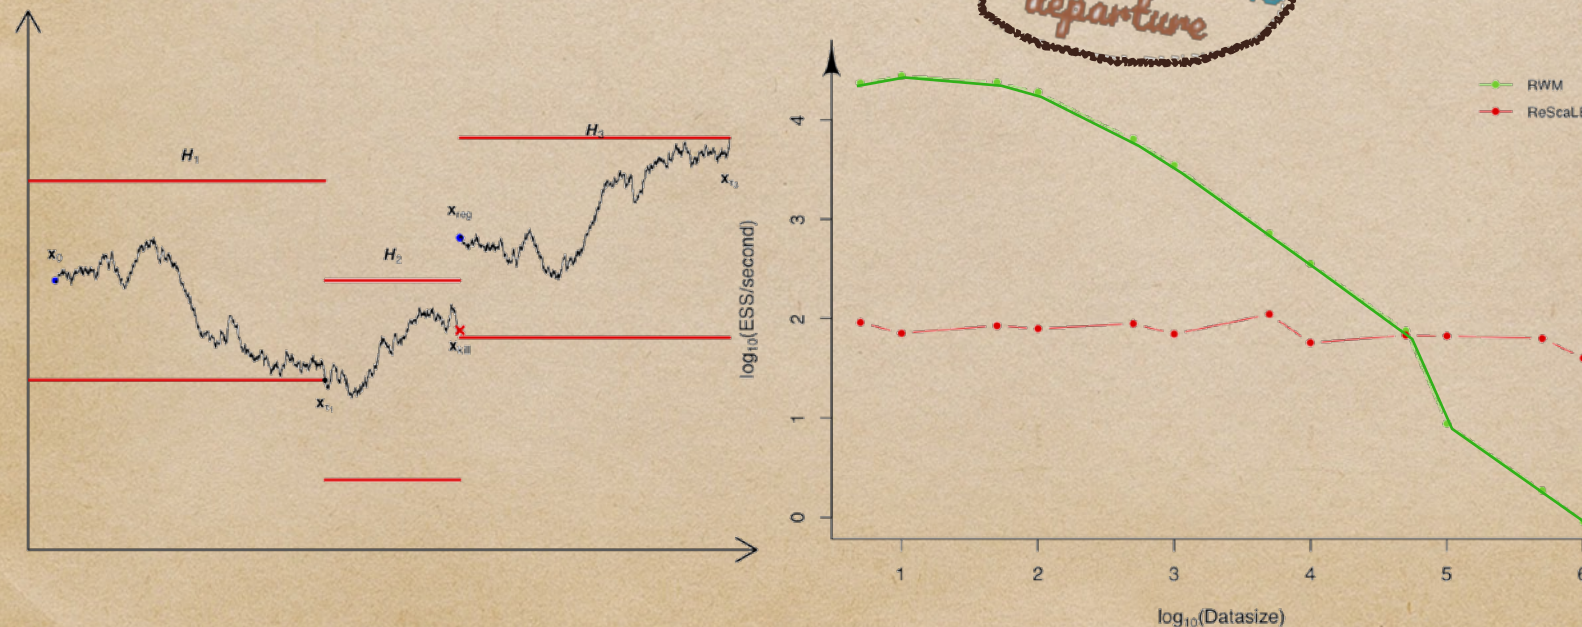


How do we deal with Big Data in Bayesian inference?



Big Data → MCMC: acceptance rate is expensive → less efficient
 (large n) → ReScaLE: hazard rate is expensive → Use subsampling → Still Converges → Efficient



By: Divakar Kumar
 Supervisors:
 Gareth Roberts &
 Murray Pollock